

Class 6: R Functions

Kira

R Functions

In this session we will work through the process of developing our own function for calculating average grades for fictional students in a fictional class.

We will start with a simplified version of the problem. Grade some vectors of student scores. We want to drop the lowest score and get the average.

```
# Example input vectors to start with
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

We can use the `mean()` function to get the average:

```
mean(student1)
```

```
[1] 98.75
```

We can find the smallest value with the `min()` function.

```
min(student1)
```

```
[1] 90
```

There is also the `which.min()` function. Let's see if this can help:

```
student1
```

```
[1] 100 100 100 100 100 100 100 90
```

```
which.min(student1)
```

```
[1] 8
```

```
student1[which.min(student1)]
```

```
[1] 90
```

```
x <- 1:5  
x
```

```
[1] 1 2 3 4 5
```

```
x[-4]
```

```
[1] 1 2 3 5
```

Let's put this together to drop the lowest value and find the average.

```
mean(student1[-which.min(student1)])
```

```
[1] 100
```

Now what about student2?

```
student2
```

```
[1] 100 NA 90 90 90 90 97 80
```

```
mean(student2[-which.min(student2)], na.rm=TRUE)
```

```
[1] 92.83333
```

Hmmm....ok what about student 3?

```
student3
```

```
[1] 90 NA NA NA NA NA NA NA
```

```
mean(student3[-which.min(student3)], na.rm=TRUE)
```

```
[1] NaN
```

```
mean(student3, na.rm=TRUE)
```

```
[1] 90
```

So this sucks! This inflates grades as it drops all the NAs before determining the mean...
It seems like the function `is.na()` might be useful.

```
student3
```

```
[1] 90 NA NA NA NA NA NA NA
```

```
is.na(student3)
```

```
[1] FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

```
student2
```

```
[1] 100 NA 90 90 90 90 97 80
```

```
is.na(student2)
```

```
[1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE
```

I can use a logical vector to index another vector.

```
student2[is.na(student2)] <- 0
student2
```

```
[1] 100    0  90  90  90  90  97  80
```

```
x <- student3
x[is.na(x)] <- 0
x
```

```
[1] 90  0  0  0  0  0  0  0
```

```
x <- student2
x[is.na(x)] <- 0
mean(x[-which.min(x)])
```

```
[1] 91
```

We have our working snippet of code! This is now going to be the body of our function.

All functions in R have at least 3 things: - A name (we pick that) - input arguments - a body (the code that does the work) #Q1: Working R function to calculate the average and drop the lowest score:

```
grade <- function(x){
  # Mask NA to zero
  x[is.na(x)] <- 0
  # Drop lowest value and get mean
  mean(x[-which.min(x)])
}
```

Make sure to press play to send it to the R “brain”. Let’s try it out with student 1.

```
grade(student1)
```

```
[1] 100
```

```
grade(student2)
```

```
[1] 91
```

```
grade(student3)
```

```
[1] 12.85714
```

```
gradebook <- read.csv("https://tinyurl.com/gradeinput",row.names=1)
head(gradebook)
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77

I can use the complicated but useful function `apply()` to use our existing `grade()` function on the whole class gradebook.

How does the `apply()` function work?

```
results <- apply(gradebook,1,grade)
results
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

Q2. Using your `grade()` function and the supplied gradebook, Who is the top scoring student overall in the gradebook? [3pts]

```
which.max(results)
```

```
student-18
18
```

Q3. From your analysis of the gradebook, which homework was toughest on students (i.e. obtained the lowest scores overall)? [2pts]

```
which.min(apply(gradebook,2,sum,na.rm=TRUE))
```

```
hw2  
2
```

Homework 2 had the lowest scores overall.

```
# not a good way  
which.min(apply(gradebook,2,mean,na.rm=TRUE))
```

```
hw3  
3
```

If I want to use the mean approach, I will need to mask the NA (missing homeworks) to zero first.

```
mask <- gradebook  
mask[is.na(mask)] <- 0  
mask
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	0	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77
student-7	89	100	74	87	100
student-8	89	100	76	86	100
student-9	86	100	77	88	77
student-10	89	72	79	0	76
student-11	82	66	78	84	100
student-12	100	70	75	92	100
student-13	89	100	76	100	80
student-14	85	100	77	89	76
student-15	85	65	76	89	0
student-16	92	100	74	89	77
student-17	88	63	100	86	78
student-18	91	0	100	87	100
student-19	91	68	75	86	79
student-20	91	68	76	88	76

```
which.min(apply(mask,2,mean,na.rm=TRUE))
```

```
hw2  
2
```

Q4. Optional Extension: From your analysis of the gradebook, which homework was most predictive of overall score (i.e. highest correlation with average grade score)? [1pt]

Here we are going to look at the correlation of each Homework results (i.e the columns in the gradebook) with the overall grade of students from the course (in the `results` object obtained from using our `grade()` function).

```
results
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

```
mask$hw4
```

```
[1] 88 89 100 100 86 89 87 86 88 0 84 92 100 89 89 89 86 87 86  
[20] 88
```

I am going to use the `cor()` function:

```
cor(results,mask$hw4)
```

```
[1] 0.3810884
```

```
cor(results,mask$hw5)
```

```
[1] 0.6325982
```

I can use the `apply()` function to find the correlation for all homework assignments at once.

```
apply(mask,2,cor, y=results)
```

hw1	hw2	hw3	hw4	hw5
0.4250204	0.1767780	0.3042561	0.3810884	0.6325982

Homework 5 was most predictive of overall score (i.e highest correlation).